

# Kubernetes

## what

- orchestration tool

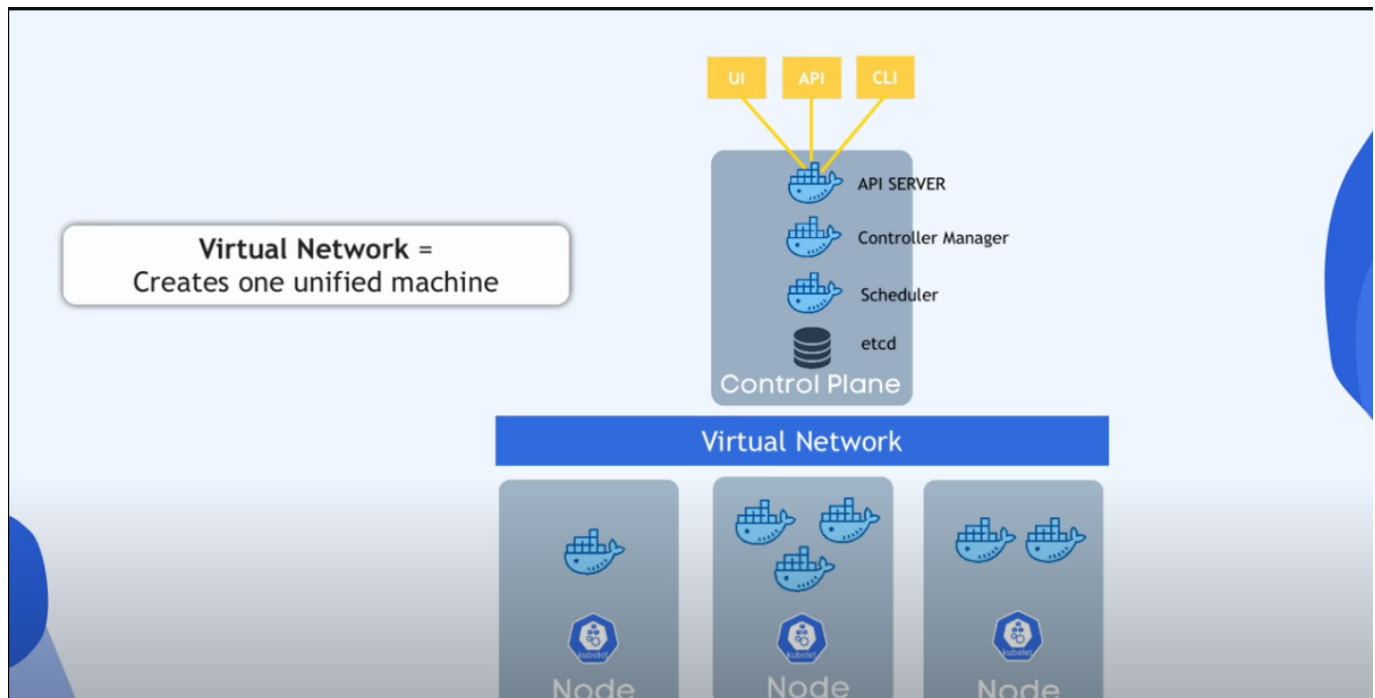
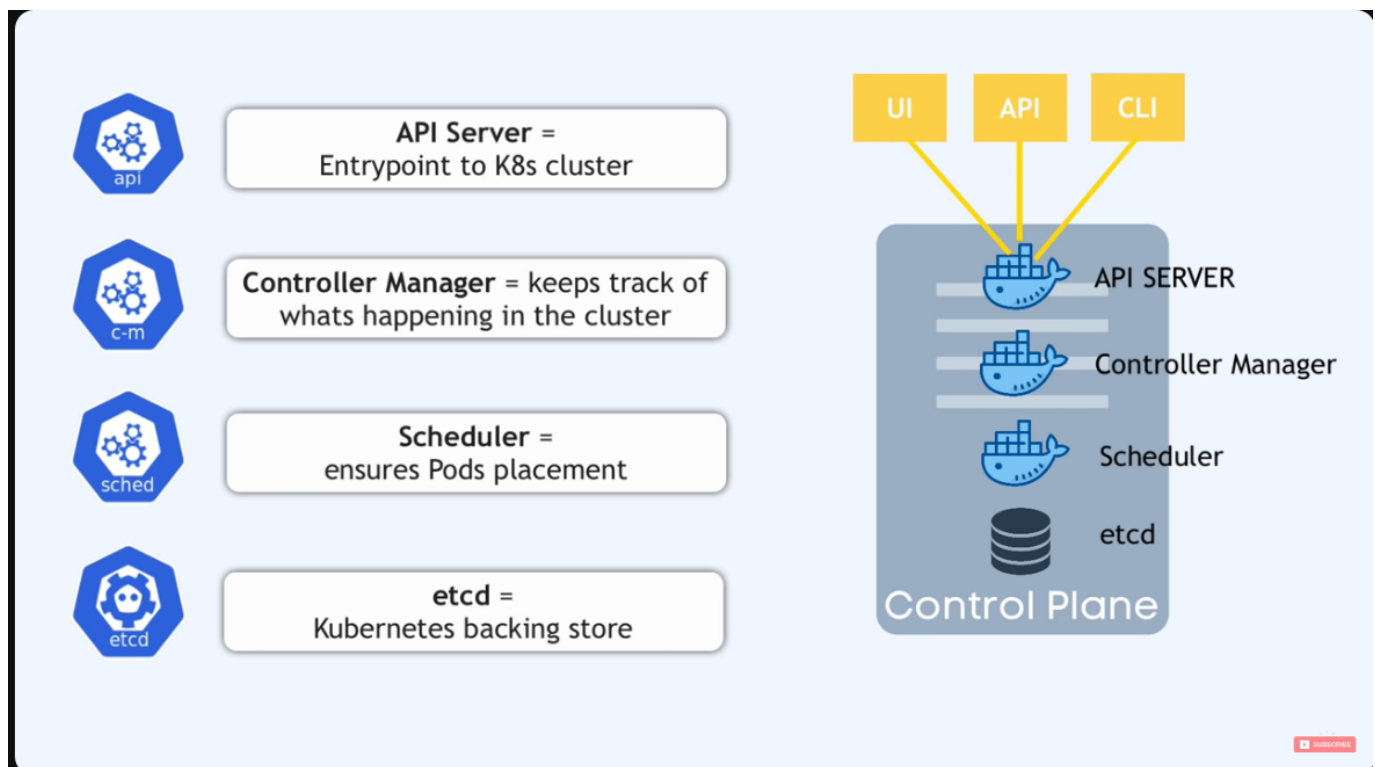
## why

- migration to microservices

## benefits

- tolerance
- availability
- scalability

## architecture



## worker nodes:

- higher workload -> much bigger resources

## master node

- handful of jobs
- very important if we lose it, we lose access to all the cluster: back it up
- duplicated most of the time

## main components

### Pod

- abstraction over container: k8s abstracts so can replace and abstracts its technology (docker or other ..)

- smallest unit
- usually one app per Pod
- one internal ip address per Pod
- ephemeral
- get assigned a new ip address when recreated (they die because of crashing for example) : problematic for changing the ip address of the pod they connect to
- that's why we need the component service

## Service

- static ip address can be attached to each Pod
- lifecycle of service and service aren't connected so we do not need to change the endpoint
- Since app should be accessible through the browser: we need an external service : opens communication from external sources
- we do not want db service accessible: we need an internal service
- internal service is the default
- can be a load balancer
- if we want to change the domain name: we use Ingress

## Ingress

- instead of service the request goes to ingress and ingress does the forwarding to service
- so far nothing fancy

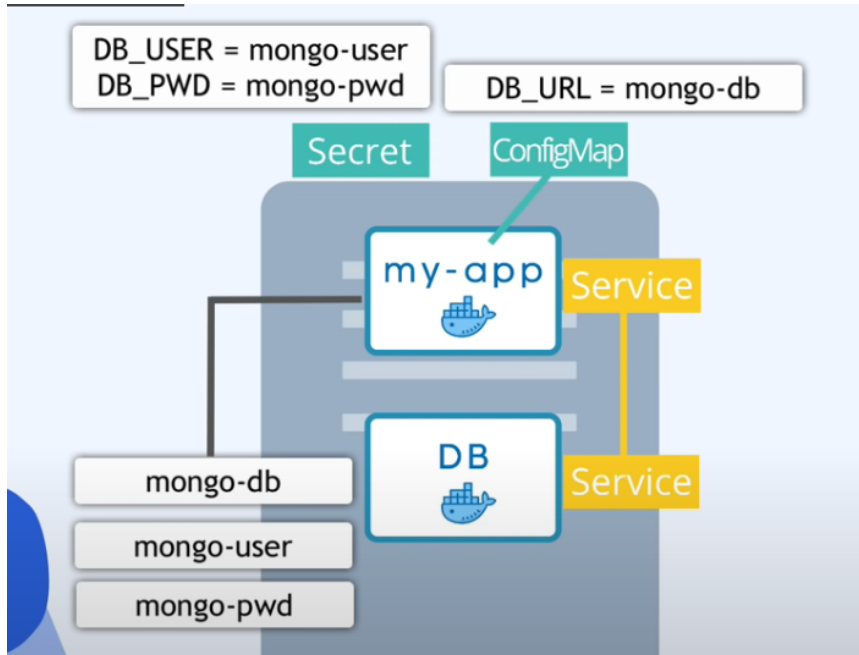
## ConfigMap

- pods communicate with Service
- DB has a URL, usually in the built application
- in case it changed: rebuild and push and pull new image to the Pod -> Tedious!
- ConfigMap: configs of the app connected to the pod
- DB user and password can change too
- ConfigMap is not for non-confidential info!

## Secret

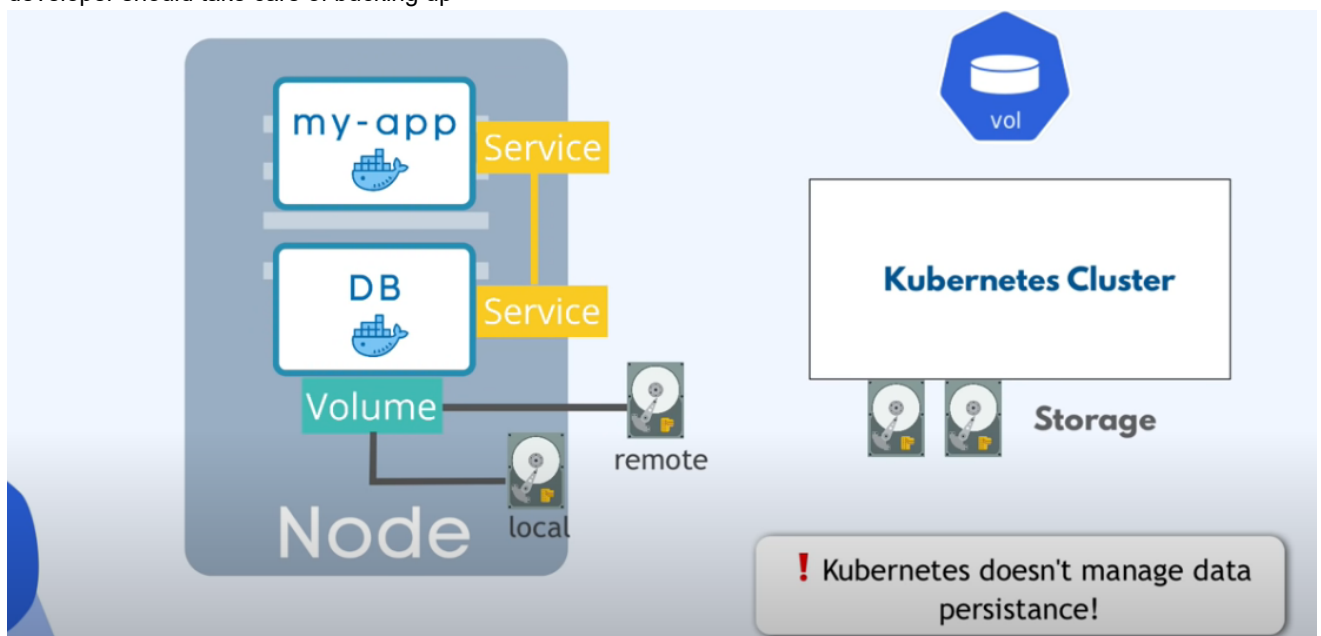
- are in base64 and not encrypted by default
- should use third party encryptor

- connect it to Pod: env variables or properties file



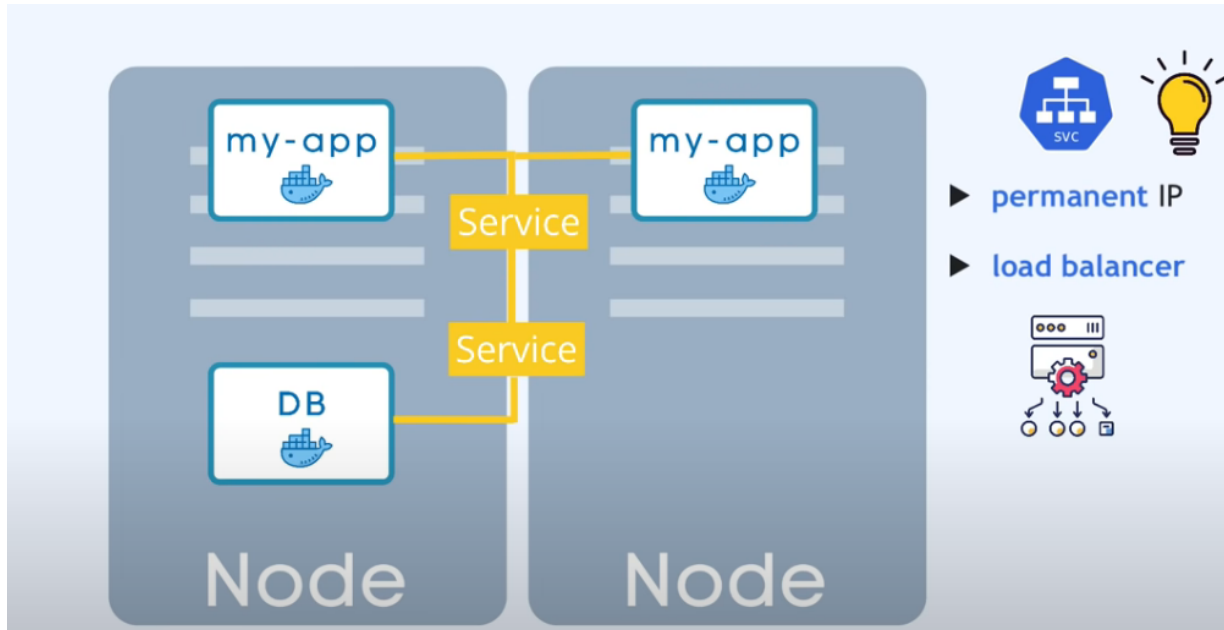
## Volume

- Pods do not persist data
- attach physical hard drive to Pod: in the local machine or outside the cluster( cloud or on premise)
- storage:external hard drive plugged into the k8s cluster
- Kubernetes doesn't manage data persistence!
- developer should take care of backing up



## Deployment

- if an application crashes, the service will forward the requests to another one




- to define blueprint for Pod and how many replicas(scale up or down)
- we mostly create deployments not pods
- deployment an abstraction over Pods
- DB can't be replicated using Deployment, because it has state
- if we want replicas, they need to access the same shared data storage to avoid data inconsistencies

## StatefulSet

- for stateful apps
- takes care of the consistency over the replicas
- using this service is tedious -> DBs are often hosted on external services

## K8s config

- goes through the master with the process API server: the only endpoint to the cluster
- requests in yaml or json
- config requests are declarative
- controller manager checks: desired state == actual state ?
- config files are with the code or in their own git repo



# Config files parts

## metadata

```
..:
kind
metadata:
  name
```

## specification

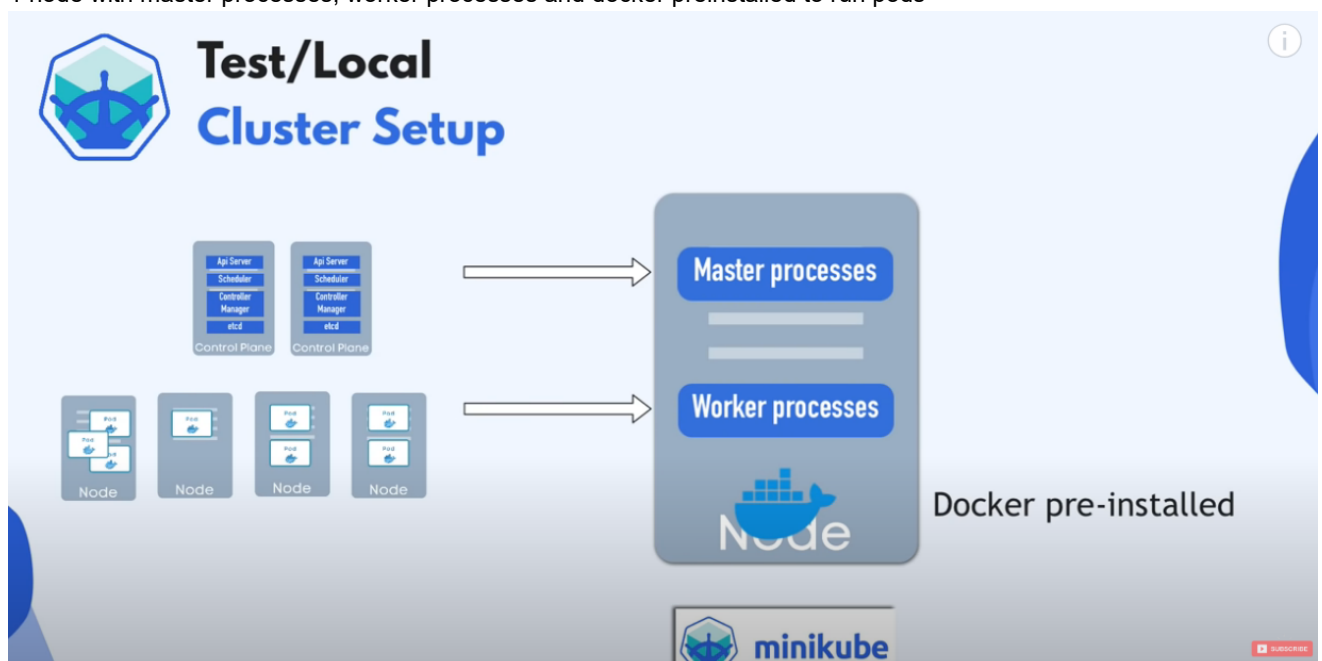
- specific to the kind

## status

- automatically generated
- checking desired == actual
- k8s updates status continuously
- the status info comes from the brain of the cluster: etcd

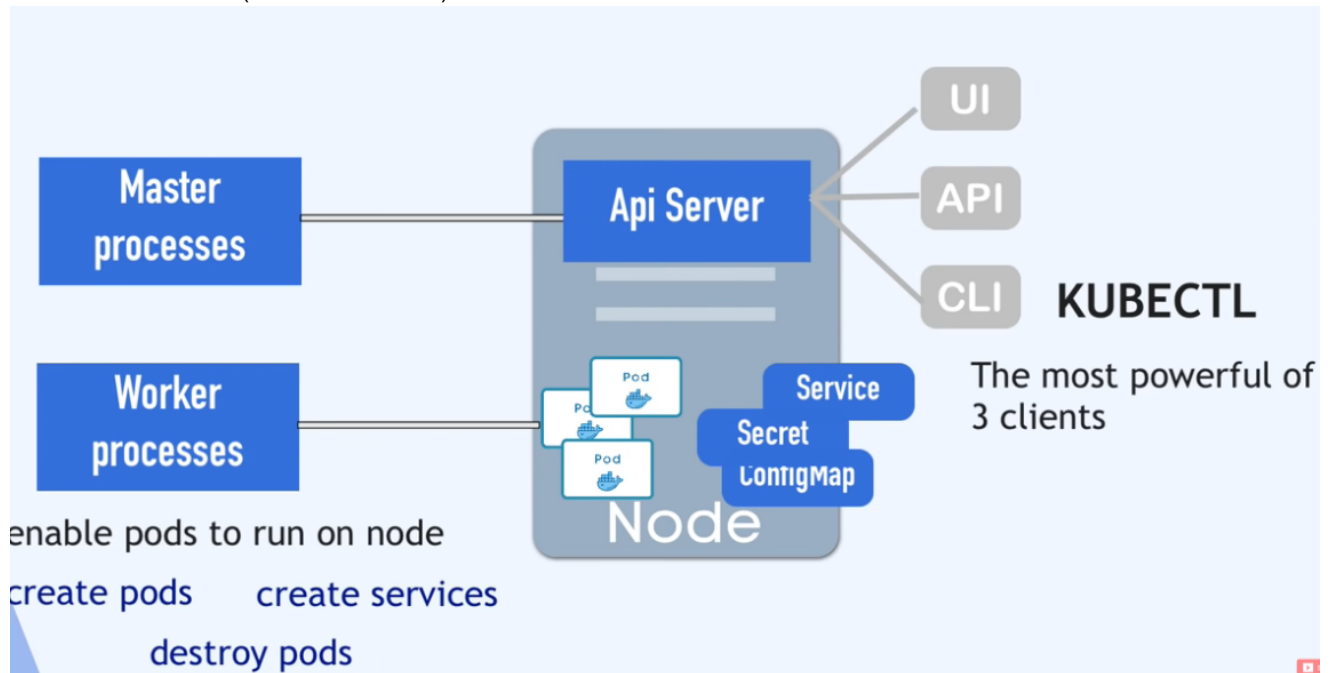
## Minikube

- if we want to test cluster set ups locally
- 1 node with master processes, worker processes and docker preinstalled to run pods



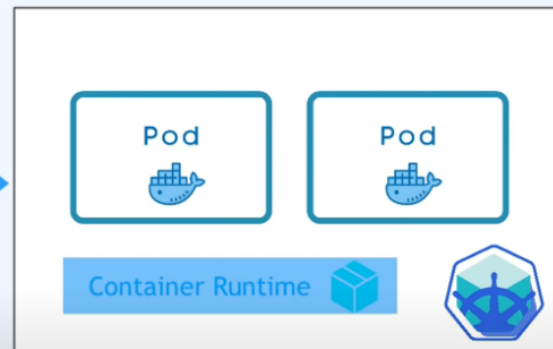
## Cubectl

- CLI to for k8s cluster (cloud or Minikube)



## 2 Layers of Docker

1) Minikube runs as Docker container



2) Docker inside Minikube to run our application containers

Minikube cluster