

# Kubernetes

## 1. Installation and Configuration of Kubernetes

- Install and configure a local or online Kubernetes cluster.
- Verify that the cluster is operational using the kubectl command.

```
PS C:\Users\salma> kubectl config current-context  
docker-desktop
```

```
kubectl config view
```

```
L$ kubectl config view  
apiVersion: v1  
clusters:  
- cluster:  
  certificate-authority-data: DATA+OMITTED  
  server: https://kubernetes.docker.internal:6443  
  name: docker-desktop  
contexts:  
- context:  
  cluster: docker-desktop  
  user: docker-desktop  
  name: docker-desktop  
current-context: docker-desktop  
kind: Config  
preferences: {}  
users:  
- name: docker-desktop  
  user:  
    client-certificate-data: DATA+OMITTED  
    client-key-data: DATA+OMITTED
```

## 2. Application Deployment

- Create a Kubernetes deployment for the web application using a YAML configuration file.

```
apiVersion: apps/v1  
  
kind: Deployment  
  
metadata:  
  
  name: first-deployment  
  
spec:  
  
  replicas: 1  
  
  selector:  
  
    matchLabels:  
  
      app: spring-app  
  
  template:  
  
    metadata:  
  
      labels:  
  
        app: spring-app
```

```
spec:

  containers:

    - name: spring-app

      image: salmaghabri/spring-app

      ports:

        - containerPort: 8080
```

PROBLEMS OUTPUT DEBUG CONSOLE **TERMINAL** PORTS COMMENTS

```
● PS C:\Users\salma\OneDrive\Bureau\9raya\DevOupsie> kubectl apply -f first_deployment.yaml
deployment.apps/first-deployment created
○ PS C:\Users\salma\OneDrive\Bureau\9raya\DevOupsie> |
```

- Verify that the deployment is done correctly and that the pods are running.

```
● PS C:\Users\salma\OneDrive\Bureau\9raya\DevOupsie> kubectl get pods --show-labels
NAME                                READY   STATUS    RESTARTS   AGE   LABELS
first-deployment-c6f7889f9-x85rx    1/1     Running   0           6m11s  app=spring-app,pod-template-hash=c6f7889f9

● PS C:\Users\salma\OneDrive\Bureau\9raya\DevOupsie> kubectl get deployments
NAME                                READY   UP-TO-DATE   AVAILABLE   AGE
first-deployment                    1/1     1             1           18m
```

### 3. Replica Management

- Scale the application by modifying the number of replicas of the deployment.

```
kubectl scale deployment first-deployment --replicas=3
```

```
first-deployment 1/1 1 1 18m
● PS C:\Users\salma\OneDrive\Bureau\9raya\DevOupsie> kubectl scale deployment first-deployment --replicas=3
deployment.apps/first-deployment scaled
```

- Verify that the new replicas are created and deployed correctly.

```
● PS C:\Users\salma\OneDrive\Bureau\9raya\DevOupsie> kubectl scale deployment first-deployment --replicas=3
deployment.apps/first-deployment scaled
● PS C:\Users\salma\OneDrive\Bureau\9raya\DevOupsie> kubectl get deployments
NAME                                READY   UP-TO-DATE   AVAILABLE   AGE
first-deployment                    3/3     3             3           21m
```

### 4. Service Definition

- Create a Kubernetes service to expose the application internally.

```
apiVersion: v1

kind: Service

metadata:

  name: spring-app-service

spec:

  selector:

    app: spring-app

  ports:
```

```
- protocol: TCP

port: 80

targetPort: 8080
```

- `selector` specifies the pods targeted by this service. It matches the pods labeled with `app: spring-app`, which is the same label used in your deployment.
- `ports` specifies the ports exposed by the service. In this example, the service exposes port 80 on the cluster, which is forwarded to port 8080 of the pods.

```
kubectl apply -f service.yaml
```

```
PS C:\Users\salma\OneDrive\Bureau\9raya\DevOpsie> kubectl apply -f first_service.yaml
service/spring-app-service created
PS C:\Users\salma\OneDrive\Bureau\9raya\DevOpsie> kubectl get services
NAME            TYPE        CLUSTER-IP    EXTERNAL-IP  PORT(S)    AGE
kubernetes      ClusterIP   10.96.0.1     <none>       443/TCP    5h17m
spring-app-service ClusterIP   10.109.155.198 <none>       80/TCP     54s
```

- Verify that the service is accessible from the Kubernetes cluster.

```
PS C:\Users\salma\OneDrive\Bureau\9raya\DevOpsie> kubectl get services
NAME            TYPE        CLUSTER-IP    EXTERNAL-IP  PORT(S)    AGE
kubernetes      ClusterIP   10.96.0.1     <none>       443/TCP    5h26m
spring-app-service ClusterIP   10.109.155.198 <none>       80/TCP     10m
PS C:\Users\salma\OneDrive\Bureau\9raya\DevOpsie> kubectl get endpoints spring-app-service
NAME            ENDPOINTS                                     AGE
spring-app-service 10.1.0.11:8080,10.1.0.12:8080,10.1.0.13:8080 13m
```

the output shows that the `spring-app-service` service is currently routing traffic to three different pods (`10.1.0.11:8080`, `10.1.0.12:8080`, and `10.1.0.13:8080`), all of which are listening on port `8080`. This suggests that our Spring application is being served by multiple pods and is ready to receive traffic.

## 5. External Application Exposure:

- Define a LoadBalancer type service to expose the application externally.

```
apiVersion: v1

kind: Service

metadata:

  name: spring-app-loadbalancer

spec:

  type: LoadBalancer

  ports:

    - protocol: TCP

      port: 80

      targetPort: 8080

  selector:

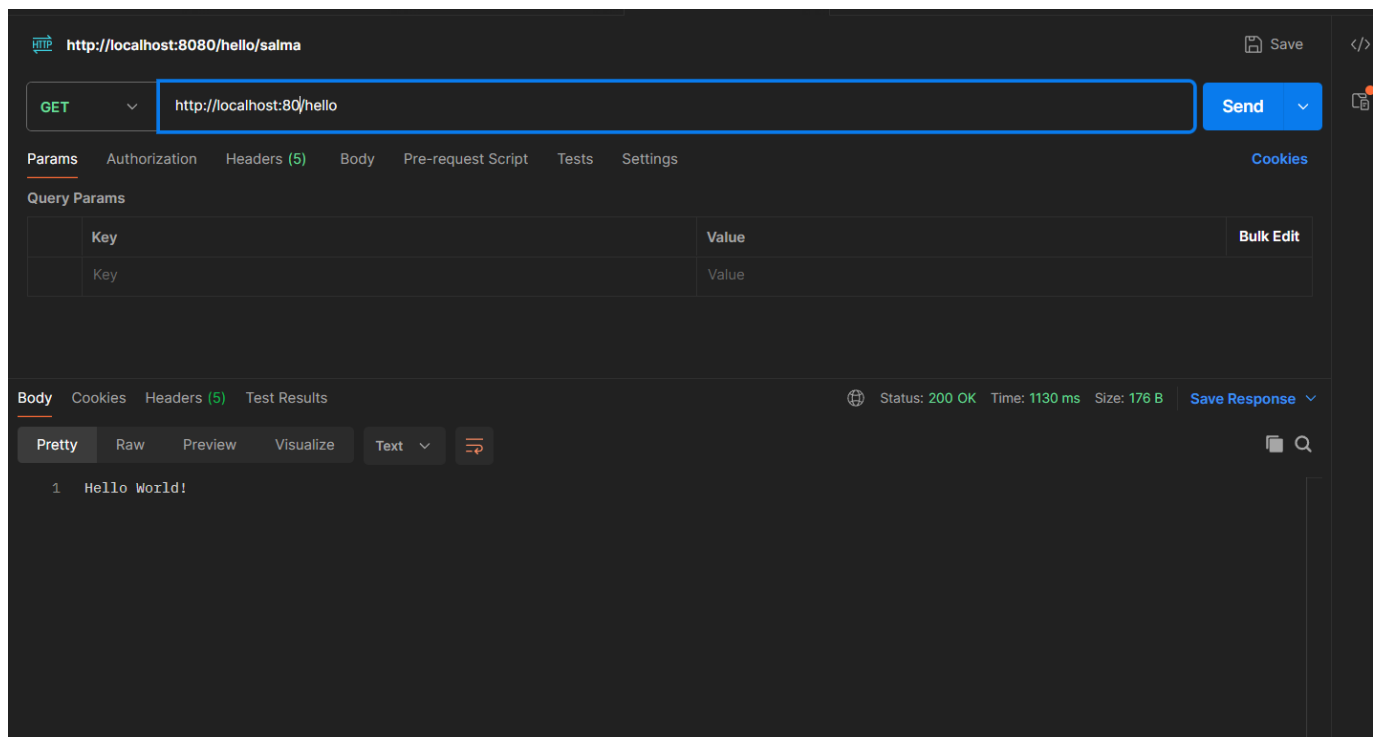
    app: spring-app
```

- Verify that the application can be accessed from outside the cluster using the assigned IP address. Once the LoadBalancer type service is created, we can check its status and get the assigned IP address. Service information, including its type, exposed ports, and the IP address assigned by the LoadBalancer.

```
kubectl get svc spring-app-loadbalancer
```

```
PS C:\Users\salma\OneDrive\Bureau\9raya\DevOupsie> kubectl apply -f load_balancer_service.yaml
service/spring-app-loadbalancer created
PS C:\Users\salma\OneDrive\Bureau\9raya\DevOupsie> kubectl get svc spring-app-loadbalancer
NAME                TYPE          CLUSTER-IP    EXTERNAL-IP  PORT(S)          AGE
spring-app-loadbalancer  LoadBalancer  10.111.25.46   localhost    80:32040/TCP     3m56s
```

On our machine we can access the app via `http://localhost:80` if we send a get request to `/hello`



## 6. Update Management

Update the container image used by the application.

We change the image to nginx for example:

```
apiVersion: apps/v1

kind: Deployment

metadata:

  name: first-deployment

spec:

  replicas: 1

  selector:

    matchLabels:

      app: spring-app
```

```

template:

  metadata:

    labels:

      app: spring-app

  spec:

    containers:

      - name: spring-app

        image: nginx // image changed

        ports:

          - containerPort: 80

```

Verify that the new version of the application is deployed correctly without service interruption.

```

PS C:\Users\salma\OneDrive\Bureau\9raya\DevOpsie> kubectl get pods
NAME                READY   STATUS    RESTARTS   AGE
debug               1/1     Running   0           90m
first-deployment-59ff986c5d-njtb5  0/1     ContainerCreating   0           28s

```

while the container is created, the old one is still listening on localhost:80

The screenshot shows a web browser interface with the URL `http://localhost:80/hello?name=ena` in the address bar. The browser's developer tools are open, showing the 'Query Params' section with a table:

Key	Value
name	ena

Below the table, the 'Body' tab is selected, showing the response: `1 Hello ena!`. The status bar at the bottom indicates 'Status: 200 OK', 'Time: 1307 ms', and 'Size: 174 B'.

## 7. Monitoring and Logging: Prometheus

[source](#)

[kube-state-metrics](#) (KSM) is a simple service that listens to the Kubernetes API server and generates metrics about the state of the objects deployed in the cluster.

Kube State Metrics is commonly used in conjunction with Prometheus, a popular open-source monitoring and alerting system.

## Install `kube-state-metrics`

```
git clone https://github.com/kubernetes/kube-state-metrics.git
```

```
kubectl apply -f kubernet
```

```
(salma@LAPTOP-C38TPMN9) - [ /mnt/c/Users/salma/OneDrive/Bureau/9raya/DevOupsie/kube-state-metrics ]
$ kubectl apply -f kubernet
clusterrolebinding.rbac.authorization.k8s.io/kube-state-metrics created
clusterrole.rbac.authorization.k8s.io/kube-state-metrics created
deployment.apps/kube-state-metrics created
serviceaccount/kube-state-metrics created
service/kube-state-metrics created
```

## Set Up `kube-state-metrics` to Expose Metrics for Your Kubernetes Cluster

We create the `kube-state-metrics` service

```
kind: Service
apiVersion: v1
metadata:
  namespace: kube-system
  name: kube-state-nodeport
spec:
  selector:
    k8s-app: kube-state-metrics
  ports:
    - protocol: TCP
      port: 8080
      nodePort: 30000
  type: NodePort
```

```
kubectl apply -f kube-state-metrics-nodeport-svc.yml
```

```
(salma@LAPTOP-C38TPMN9) - [ /mnt/c/Users/salma/OneDrive/Bureau/9raya/DevOupsie ]
$ kubectl apply -f kube-state-metrics-nodeport-svc.yml
service/kube-state-nodeport created
```

If we open the browser and go to `localhost:30000` we can see the data of our cluster

```
# my global config

global:

    scrape_interval: 15s # Set the scrape interval to every 15 seconds. Default is every 1 minute.
    evaluation_interval: 15s # Evaluate rules every 15 seconds. The default is every 1 minute.

    # scrape_timeout is set to the global default (10s).

# Alertmanager configuration
```

```

alerting:

  alertmanagers:

    - static_configs:

      - targets:

        # - alertmanager:9093

# Load rules once and periodically evaluate them according to the global 'evaluation_interval'.

rule_files:

# - "first_rules.yml"

# - "second_rules.yml"

# A scrape configuration containing exactly one endpoint to scrape:

# Here it's Prometheus itself.

scrape_configs:

# The job name is added as a label 'job=<job_name>' to any timeseries scraped from this config.

- job_name: "prometheus"

# metrics_path defaults to '/metrics'

# scheme defaults to 'http'.

static_configs:

  - targets: ["localhost:9090"]

- job_name: kubernetes


static_configs:

  - targets: ["localhost:30000"]

```

## Example of a query

```
kube_pod_status_ready{namespace="default",condition="true"}
```


Prometheus
Alerts
Graph
Status ▾
Help

☐ Use local time
☐ Enable query history
☒ Enable autocomplete
☒ Enable highlighting
☒ Enable linter

⌕
⌵
⌶
Execute

Table
Graph

Load time: 62ms
Resolution: 14s
Result series: 2

<
>
Evaluation time

kube_pod_status_ready(condition="true",instance="localhost:30000",job="kubernetes",namespace="default",pod="debug")	0
kube_pod_status_ready(condition="true",instance="localhost:30000",job="kubernetes",namespace="default",pod="first-deployment-c6f7889f9-q628s")	1

Remove Panel

Add Panel



